

Yoav Gur Arie

☎ 054-941-2753 ✉ yoavgurarieh@gmail.com in @yoav-gur-arie

Graduate researcher in NLP and AI with 9 years of software engineering experience, specializing in interpretability and knowledge adaptation in large language models.

EDUCATION

Tel Aviv University — Computer Science MSc Student, *The Adi Lautman Interdisciplinary Program for Outstanding Students*

2022 – 2025

- Third-year student at Tel Aviv University's Lautman honors program, pursuing a computer science degree within a four-year **direct-to-masters** track emphasizing interdisciplinary study, with a focus on computer science.
- Engaging in diverse coursework spanning mathematics, computational biology, economics, and history, fostering a broad skill set and interdisciplinary understanding.
- Taking advantage of the program's flexibility in order to take high-level courses regardless of prerequisites, independently meeting their demands through self-learning.

Current grade average: **93.52**

EXPERIENCE

Graduate Researcher – NLP and Interpretability (Dr. Mor Geva's Lab)

August 2024 – Present

- Researching novel methods of knowledge adaptation in LLMs, employing interpretability techniques to analyze inter-module communication and optimize knowledge integration within the model.
- Working on new methods for large scale automatic interpretation of sparse auto-encoder features, emphasizing compute and memory efficiency, while defining rigorous criteria for effective explanations.
- Establishing a large-scale data collection pipeline with hermetically isolated and parsed entities and facts, designed for training LLMs with precise knowledge attribution.

Laminar Security / Rubrik — Senior Software Developer

March 2021 - Present

- Joined the startup during its inception and spearheaded the development of crucial front-end, back-end, and agent-based systems from scratch.
- Led the development of a document sensitivity classification engine, involving parsing documents, analyzing their content, and categorizing their sensitivity levels.
- Conducted technical research on extracting decrypted data from encrypted traffic in cloud-based clusters, and successfully implemented an efficient solution.

IDF - 8200 Unit (Subunit 7149) — Senior Developer & Researcher

March 2016 - March 2021

- Carried out months-long solo research projects into esoteric and opaque technologies that have little or no publicly available documentation, culminating in the development of technical solutions.
- Led through to completion highly complex, high-risk projects involving multiple teams.
- Worked on smaller sporadic projects entailing deep dives into complex topics.

Checkpoint Software Technologies — Software Developer

August 2015 - March 2016

- Contributed to the development of a log aggregator and analyzer product tailored for large enterprises. Designed and implemented automated reports featuring visualization of insights.

SKILLS

Python • PyTorch • HF
Transformers • TransformerLens •
SAE Lens • NLP Interpretability •
Golang • C • Javascript • CSS • eBPF
• Linux Internals • Embedded
systems • Networks • Cyber
Security • Reverse Engineering

AWARDS

Personal award from the Military Intelligence Directorate for excellence (including a rank promotion)

Award for outstanding project (8200)

Worked on two separate projects that received the Israel Defense Prize

First place team in the 2015 'Skillz' coding competition

PROJECTS

Avian Neuronal Response Project
– Analyzed avian neuronal activity in response to varied bird calls (authentic and artificial), and used machine learning techniques to classify stimuli.

Parkalot – Developed an app and backend infrastructure to display current space availability in Tel Aviv parking lots and predict future trends.

Hoppa – Developed a platformer genre android game using the Unity game engine.

LANGUAGES

Hebrew and English – Native proficiency.

SCORES

Psychometric: 750
Bagrut Average: 112