

Yoav Gur-Arieh

PhD Student researching interpretability in LLMs, with 11 years software engineering experience

📍 Tel-Aviv, Israel ✉ yoavgurarieh@gmail.com 🔗 yoav.ml in yoavgurarieh 🌐 yoavgur

Education

Tel-Aviv University Dec 2025 to present
PhD in Computer Science

Tel-Aviv University Oct 2024 to Dec 2025
MS in Computer Science, summa cum laude (GPA 98/100)

- Done as part of the direct-to-masters interdisciplinary program.

Tel-Aviv University Oct 2022 to Jun 2026
Interdisciplinary Studies (GPA 94/100)

- Studying in the Adi Lautman Interdisciplinary Program for Outstanding Students.
- Took courses in computer science, biology, neuroscience, physics and history.

Fellowships & Scholarships

VATAT Doctoral Fellowship in AI & Data Science 2026 to 2029

- National three-year fellowship from the Israeli Council for Higher Education's Planning and Budgeting Committee (VATAT); one of up to 20 awarded annually across all universities.

TAD Excellence Fellowship for Doctoral Students 2026 to 2029

- Three-year excellence fellowship awarded by Tel Aviv University's Center for AI & Data Science (TAD).

Adi Lautman Interdisciplinary Program for Outstanding Students 2022 to 2026

- Full tuition scholarship throughout undergraduate and graduate studies.

Experience

Graduate Researcher (Advisor: [Mor Geva](#)) Sep 2024 to present
Tel-Aviv University

- Published research on improvements to large-scale automated **feature interpretability pipelines**, and **machine unlearning**.
- Working on understanding and controlling **reasoning** in LLMs.

AI Researcher Feb 2026 to present
AAI

- Working on designing expert AI systems.

Research Fellow Jun 2026 to Aug 2026
Cambridge Boston Alignment Initiative (CBAI)

Research Intern (Supervisor: [Atticus Geiger](#)) Jun to Sep 2025
Pr(Ai)²R Group

Senior Software Engineer Mar 2021 to Feb 2026
Laminar Security / Rubrik

- Joined the startup at its inception, spearheading the development of front-end, back-end, and agent-based systems from scratch, including a document sensitivity classification engine.
- Led technical research and implemented solutions for extracting decrypted data from encrypted cloud traffic.

Senior Developer & Researcher Mar 2016 to Mar 2021
Stealth

- Carried out months-long solo research projects into esoteric and opaque technologies that have little or no publicly available documentation, culminating in the development of technical solutions.

- Led through to completion highly complex, high-risk projects involving multiple teams.

Software Developer

Aug 2015 to Mar 2016

Checkpoint Software Technologies

- Contributed to the development of a log aggregator and analyzer product tailored for large enterprises.

Publications

Faithfulness Metrics Don't Measure Faithfulness: A Meta-Evaluation with Ground Truth	May 2026
Yoav Gur-Arieh, Ana Marasović, Mor Geva	
Submitted to NeurIPS 2026	
Oral presentation at NEMI 2026	
Disentangling MLP Neuron Weights in Vocabulary Space	Apr 2026
Asaf Avrahamy, Yoav Gur-Arieh, Mor Geva	
Accepted to COLM 2026	
Mixing Mechanisms: How Language Models Retrieve Bound Entities In-Context	Oct 2025
Yoav Gur-Arieh, Mor Geva, Atticus Geiger	
Accepted to ICLR 2026	
LMEnt: A Suite for Analyzing Knowledge in Language Models from Pretraining Data to Representations	Sep 2025
Daniela Gottesman, Alon Gilae-Dotan, Ido Cohen, Yoav Gur-Arieh, Marius Mosbach, Ori Yoran, Mor Geva	
Accepted to TACL	
Precise In-Parameter Concept Erasure in Large Language Models	May 2025
Yoav Gur-Arieh, Clara Suslik, Yihuai Hong, Fazl Barez, Mor Geva	
Accepted to EMNLP Main 2025	
Enhancing Automated Interpretability with Output-Centric Feature Descriptions	Jan 2025
Yoav Gur-Arieh, Roy Mayan, Chen Agassy, Atticus Geiger, Mor Geva	
Accepted to ACL Main 2025	

Invited Talks

Technion, NLP Group – <i>Evaluating Chain-of-Thought Faithfulness in LLMs</i>	Aug 2026
Transluce – <i>Mixing Mechanisms: How LMs Use Pointers to Bind and Retrieve Entities</i>	Aug 2026
Harvard University, Kempner Institute PRISM Group – <i>Evaluating Chain-of-Thought Faithfulness in LLMs</i>	Aug 2026
Northeastern University, Bau Lab – <i>Evaluating Chain-of-Thought Faithfulness in LLMs</i>	Jul 2026
Hebrew University, NLP Group – <i>Mixing Mechanisms: How LMs Use Pointers to Bind and Retrieve Entities</i>	Jan 2026
Max Planck Institute for Security and Privacy – <i>From Description to Erasure: Feature-Based Control of LLMs</i>	Sep 2025

Technologies

Languages: Python, C, Golang, C#, Java, JavaScript

Technologies: PyTorch, HF Transformers, TransformerLens, SAELens, eBPF, Linux Internals, Cybersecurity

Language Proficiency

English, Hebrew: Native