

Yoav Gur-Arieh

PhD Student researching interpretability in LLMs, with 11 years software engineering experience

📍 Tel-Aviv, Israel ✉ yoavgurarieh@gmail.com 🔗 yoav.ml in yoavgurarieh 🌐 yoavgur

Education

Tel-Aviv University

Dec 2025 to present

PhD in Computer Science

- NLP lab, researching interpretability in LLMs under advisor Mor Geva.
- Working on understanding and controlling reasoning processes in LLMs.

Tel-Aviv University

Oct 2024 to Dec 2025

MS in Computer Science — *summa cum laude*

- NLP lab, researching interpretability in LLMs under advisor Mor Geva.
- Done as part of the direct-to-masters interdisciplinary program.
- Average: 98

Tel-Aviv University

Oct 2022 to Jun 2026

Interdisciplinary Studies

- Studying in the Adi Lautman Interdisciplinary Program for Outstanding Students.
- Took courses in computer science, biology, neuroscience, physics and history.
- Average: 94

Experience

Graduate Researcher (*Advisor: Mor Geva*)

Sep 2024 to present

Tel-Aviv University

- Published research on improvements to large-scale automated **feature interpretability pipelines**, and **machine unlearning**.
- Working on understanding and controlling **reasoning** in LLMs.

AI Researcher

Feb 2026 to present

AAI

- Working on designing expert AI systems.

Research Fellow

Jun 2026 to Aug 2026

Cambridge Boston Alignment Initiative (CBAI)

Research Intern (*Supervisor: Atticus Geiger*)

Jun to Sep 2025

Pr(Ai)²R Group

Senior Software Engineer

Mar 2021 to Feb 2026

Laminar Security / Rubrik

- Joined the startup at its inception, spearheading the development of front-end, back-end, and agent-based systems from scratch, including a document sensitivity classification engine.
- Led technical research and implemented solutions for extracting decrypted data from encrypted cloud traffic.

Senior Developer & Researcher

Mar 2016 to Mar 2021

Stealth

- Carried out months-long solo research projects into esoteric and opaque technologies that have little or no publicly available documentation, culminating in the development of technical solutions.
- Led through to completion highly complex, high-risk projects involving multiple teams.

Software Developer

Aug 2015 to Mar 2016

Checkpoint Software Technologies

- Contributed to the development of a log aggregator and analyzer product tailored for large enterprises.

Publications

Faithfulness Metrics Don't Measure Faithfulness: A Meta-Evaluation with Ground Truth *May 2026*

Yoav Gur-Arieh, Ana Marasović, Mor Geva

[Submitted](#) to NeurIPS 2026

Oral presentation at [NEMI 2026](#)

Disentangling MLP Neuron Weights in Vocabulary Space *Apr 2026*

Asaf Avrahamy, **Yoav Gur-Arieh**, Mor Geva

Accepted to [COLM](#)

Mixing Mechanisms: How Language Models Retrieve Bound Entities In-Context *Oct 2025*

Yoav Gur-Arieh, Mor Geva, Atticus Geiger

Accepted to [ICLR 2026](#)

LMEnt: A Suite for Analyzing Knowledge in Language Models from Pretraining Data to Representations *Sep 2025*

Daniela Gottesman, Alon Gilae-Dotan, Ido Cohen, **Yoav Gur-Arieh**, Marius Mosbach, Ori Yoran, Mor Geva

Accepted to [TACL](#)

Precise In-Parameter Concept Erasure in Large Language Models *May 2025*

Yoav Gur-Arieh, Clara Suslik, Yihuai Hong, Fazl Barez, Mor Geva

Accepted to [EMNLP Main 2025](#)

Enhancing Automated Interpretability with Output-Centric Feature Descriptions *Jan 2025*

Yoav Gur-Arieh, Roy Mayan, Chen Agassy, Atticus Geiger, Mor Geva

Accepted to [ACL Main 2025](#)

Invited Talks

Technion, NLP Group – *Evaluating Chain-of-Thought Faithfulness in LLMs* *Aug 2026*

Transluce – *Mixing Mechanisms: How LMs Use Pointers to Bind and Retrieve Entities* *Aug 2026*

Harvard University, Kempner Institute PRISM Group – *Evaluating Chain-of-Thought Faithfulness in LLMs* *Aug 2026*

Northeastern University, Bau Lab – *Evaluating Chain-of-Thought Faithfulness in LLMs* *Jul 2026*

Hebrew University, NLP Group – *Mixing Mechanisms: How LMs Use Pointers to Bind and Retrieve Entities* *Jan 2026*

Max Planck Institute for Security and Privacy – *From Description to Erasure: Feature-Based Control of LLMs* *Sep 2025*

Projects

SAE Knowledge Erasure Project [Code](#) 

- Leveraged output-centric feature descriptions to identify MLP SAE features associated with specific concepts in the Gemma-2 2B model. Ablated these features to effectively erase the corresponding concepts, demonstrating targeted knowledge manipulation - explained in detail [here](#).

Avian Neuronal Response Project [Code](#) 

- Analyzed avian neuronal activity in response to varied bird calls (authentic and artificial), and used classical machine learning techniques to classify stimuli.

Parkalot - Parking Finder App [Link](#) 

- Developed an app that displays parking lots with free spots around the user.

Hoppa - Platform Jumper Android Game [Link](#) 

- Developed an android platform jumper game using Unity in C#, from conception to deployment.

Technologies

Languages: Python, C, Golang, C#, Java, JavaScript

Technologies: PyTorch, HF Transformers, TransformerLens, SAELens, eBPF, Linux Internals, Cybersecurity

Language Proficiency

English, Hebrew: Native