

Yoav Gur-Arieh

PhD Student and software engineer with 11 years experience, researching interpretability in LLMs

📍 Tel-Aviv, Israel 📩 yoavgurarieh@gmail.com 🌐 yoav.ml 💬 yoavgurarieh 📱 yoavgur

Education

Tel-Aviv University

PhD in Computer Science

Dec 2025 to present

- NLP lab, researching interpretability in LLMs under advisor Mor Geva.
- Working on understanding and controlling reasoning processes in LLMs.

Tel-Aviv University

MS in Computer Science

Oct 2024 to Dec 2025

- NLP lab, researching interpretability in LLMs under advisor Mor Geva.
- Done as part of the direct-to-masters interdisciplinary program.
- Average: 98

Tel-Aviv University

Interdisciplinary Studies

Oct 2022 to Jun 2026

- Studying in the Adi Lautman Interdisciplinary Program for Outstanding Students.
- Took courses in computer science, biology, neuroscience, physics and history.
- Average: 94

Experience

Graduate Researcher (*Advisor: Mor Geva*)

Tel-Aviv University

Sep 2024 to present

- Working on research in **interpretability** in LLMs
- Published research on improvements to large-scale automated **feature interpretability pipelines**, and **machine unlearning**.
- Working on understanding and controlling **reasoning** in LLMs.

AI Researcher

AAI

Feb 2026 to present

- Working on designing expert AI systems.

Research Intern (*Supervisor: Atticus Geiger*)

Pr(Ai)²R Group

Jun to Sep 2025

- Used **causal abstraction** to uncover how LLMs perform entity binding and retrieval.

Senior Software Engineer

Laminar Security / Rubrik

Mar 2021 to Feb 2026

- Joined the startup at its inception, spearheading the development of front-end, back-end, and agent-based systems from scratch, including a document sensitivity classification engine.
- Led technical research and implemented solutions for extracting decrypted data from encrypted cloud traffic.

Senior Developer & Researcher

Stealth

Mar 2016 to Mar 2021

- Carried out months-long solo research projects into esoteric and opaque technologies that have little or no publicly available documentation, culminating in the development of technical solutions.
- Led through to completion highly complex, high-risk projects involving multiple teams.

Software Developer

Checkpoint Software Technologies

Aug 2015 to Mar 2016

- Contributed to the development of a log aggregator and analyzer product tailored for large enterprises.

Publications

Mixing Mechanisms: How Language Models Retrieve Bound Entities In-Context

October 2025

Yoav Gur-Arieh, Mor Geva, Atticus Geiger

Accepted to [ICLR Main 2026](#)

LMEnt: A Suite for Analyzing Knowledge in Language Models from Pre-training Data to Representations

September 2025

Daniela Gottesman, Alon Gilae-Dotan, Ido Cohen, **Yoav Gur-Arieh**, Marius Mosbach, Ori Yoran, Mor Geva
[Submitted](#) to TACL

Precise In-Parameter Concept Erasure in Large Language Models

May 2025

Yoav Gur-Arieh, Clara Suslik, Yihuai Hong, Fazl Barez, Mor Geva

Accepted to [EMNLP Main 2025](#)

Enhancing Automated Interpretability with Output-Centric Feature Descriptions

Jan 2025

Yoav Gur-Arieh, Roy Mayan, Chen Agassy, Atticus Geiger, Mor Geva

Accepted to [ACL Main 2025](#)

Invited Talks

Hebrew University NLP Group – Mixing Mechanisms: How LMs Use Pointers to Bind and Retrieve Entities

Jan 2026

Max Planck Institute for Security and Privacy – From Description to Erasure: Feature-Based Control of LLMs

Sep 2025

Projects

SAE Knowledge Erasure Project

[Code ↗](#)

- Leveraged output-centric feature descriptions to identify MLP SAE features associated with specific concepts in the Gemma-2 2B model. Ablated these features to effectively erase the corresponding concepts, demonstrating targeted knowledge manipulation - explained in detail [here](#).

Avian Neuronal Response Project

[Code ↗](#)

- Analyzed avian neuronal activity in response to varied bird calls (authentic and artificial), and used classical machine learning techniques to classify stimuli.

Parkalot - Parking Finder App

[Link ↗](#)

- Developed an app that displays parking lots with free spots around the user.
- Created the scraper (python), backend (python), database (SQL) and app (JS).

Hoppa - Platform Jumper Android Game

[Link ↗](#)

- Developed an android platform jumper game using Unity in C#, from conception to deployment.

Technologies

Languages: Python, C, Golang, C#, Java, JavaScript

Technologies: PyTorch, HF Transformers, TransformerLens, SAELens, eBPF, Linux Internals, Cybersecurity

Language Proficiency

English, Hebrew: Native